

# Ridge and Lasso impemenattion on Advertising dataset

In [1]:

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.linear_model import Ridge, RidgeCV, Lasso
from sklearn.preprocessing import StandardScaler
```

In [2]:

```
data=pd.read_csv(r"C:\Users\chila\Downloads\Advertising.csv")
data
```

Out[2]:

	TV	Radio	Newspaper	Sales
0	230.1	37.8	69.2	22.1
1	44.5	39.3	45.1	10.4
2	17.2	45.9	69.3	12.0
3	151.5	41.3	58.5	16.5
4	180.8	10.8	58.4	17.9
...	...	...	...	...
195	38.2	3.7	13.8	7.6
196	94.2	4.9	8.1	14.0
197	177.0	9.3	6.4	14.8
198	283.6	42.0	66.2	25.5
199	232.1	8.6	8.7	18.4

200 rows × 4 columns

In [3]:

```
data.head
```

Out[3]:

```
<bound method NDFrame.head of
0    230.1    37.8    69.2    22.1
1     44.5    39.3    45.1    10.4
2     17.2    45.9    69.3    12.0
3    151.5    41.3    58.5    16.5
4    180.8    10.8    58.4    17.9
..     ...     ...     ...     ...
195    38.2     3.7    13.8     7.6
196    94.2     4.9     8.1    14.0
197   177.0     9.3     6.4    14.8
198   283.6    42.0    66.2    25.5
199   232.1     8.6     8.7    18.4
```

[200 rows x 4 columns]>

In [4]:

```
data.tail()
```

Out[4]:


	TV	Radio	Newspaper	Sales
195	38.2	3.7	13.8	7.6
196	94.2	4.9	8.1	14.0
197	177.0	9.3	6.4	14.8
198	283.6	42.0	66.2	25.5
199	232.1	8.6	8.7	18.4

In [5]:

```
plt.figure(figsize = (10, 10))
sns.heatmap(data.corr(), annot = True)
```


Out[5]:

<Axes: >



In [6]:

```
data.drop(columns = ["Radio", "Newspaper"], inplace = True)
#pairplot
sns.pairplot(data)
data.Sales = np.log(data.Sales)
```



In [7]:

```
features = data.columns[0:2]
target = data.columns[-1]
#X and y values
X = data[features].values
y = data[target].values
#split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=17)
print("The dimension of X_train is {}".format(X_train.shape))
print("The dimension of X_test is {}".format(X_test.shape))
#Scale features
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)
```

The dimension of X\_train is (140, 2)  
The dimension of X\_test is (60, 2)

In [8]:

```
#Model
lr = LinearRegression()
#Fit model
lr.fit(X_train, y_train)
#predict
#prediction = lr.predict(X_test)
#actual
actual = y_test
train_score_lr = lr.score(X_train, y_train)
test_score_lr = lr.score(X_test, y_test)
print("\nLinear Regression Model:\n")
print("The train score for lr model is {}".format(train_score_lr))
print("The test score for lr model is {}".format(test_score_lr))
```

Linear Regression Model:

The train score for lr model is 1.0  
The test score for lr model is 1.0

In [9]:


```
#Ridge Regression Model
ridgeReg = Ridge(alpha=10)
ridgeReg.fit(X_train,y_train)
#train and test score for ridge regression
train_score_ridge = ridgeReg.score(X_train, y_train)
test_score_ridge = ridgeReg.score(X_test, y_test)
print("\nRidge Model:\n")
print("The train score for ridge model is {}".format(train_score_ridge))
print("The test score for ridge model is {}".format(test_score_ridge))
```

Ridge Model:

The train score for ridge model is 0.9902871391941609  
The test score for ridge model is 0.984426628514122

In [10]:

```
plt.figure(figsize = (10, 10))
plt.plot(features,ridgeReg.coef_,alpha=0.7,linestyle='none',marker='*',markersize=5,color='red',label=r'Ridge; $\alpha = 10$')
#plt.plot(rr100.coef_,alpha=0.5,linestyle='none',marker='d',markersize=6,color='blue',label=r'Ridge; $\alpha = 100$')
plt.plot(features,lr.coef_,alpha=0.4,linestyle='none',marker='o',markersize=7,color='green',label='Linear Regression')
plt.xticks(rotation = 90)
plt.legend()
plt.show()
```



In [11]:

```
#Lasso regression model
print("\nLasso Model: \n")
lasso = Lasso(alpha = 10)
lasso.fit(X_train,y_train)
train_score_ls =lasso.score(X_train,y_train)
test_score_ls =lasso.score(X_test,y_test)
print("The train score for ls model is {}".format(train_score_ls))
print("The test score for ls model is {}".format(test_score_ls))
```

Lasso Model:

The train score for ls model is 0.0


The test score for ls model is -0.0042092253233847465

In [12]:

```
pd.Series(lasso.coef_, features).sort_values(ascending = True).plot(kind = "bar")
```

Out[12]:

<Axes: >



In [13]:


```
#Using the linear CV model
from sklearn.linear_model import LassoCV
#Lasso Cross validation
lasso_cv = LassoCV(alphas = [0.0001, 0.001,0.01, 0.1, 1, 10], random_state=0).fit(X_train, y_train)
#score
print(lasso_cv.score(X_train, y_train))
print(lasso_cv.score(X_test, y_test))
```

0.9999999343798134

0.9999999152638072

In [14]:

```
#plot size
plt.figure(figsize = (10, 10))
#add plot for ridge regression
plt.plot(features,ridgeReg.coef_,alpha=0.7,linestyle='none',marker='*',markersize=5,color='red',label=r'Ridge;  $\alpha=0.7$ ')
#add plot for lasso regression
plt.plot(lasso_cv.coef_,alpha=0.5,linestyle='none',marker='d',markersize=6,color='blue',label=r'lasso;  $\alpha = \text{grid}$ ')
#add plot for linear model
plt.plot(features,lr.coef_,alpha=0.4,linestyle='none',marker='o',markersize=7,color='green',label='Linear Regression')
#rotate axis
plt.xticks(rotation = 90)
plt.legend()
plt.title("Comparison plot of Ridge, Lasso and Linear regression model")
plt.show()
```



In [15]:

```
#Using the Linear CV model
from sklearn.linear_model import RidgeCV
#Ridge Cross validation
ridge_cv = RidgeCV(alphas = [0.0001, 0.001, 0.01, 0.1, 1, 10]).fit(X_train, y_train)
#score
print("The train score for ridge model is {}".format(ridge_cv.score(X_train, y_train)))
print("The train score for ridge model is {}".format(ridge_cv.score(X_test, y_test)))
```

The train score for ridge model is 0.99999999997627  
The train score for ridge model is 0.999999999962466

# Ridge and Lasso impemenattion on Vehicals dataset

In [16]:

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.linear_model import Ridge, RidgeCV, Lasso
from sklearn.preprocessing import StandardScaler
```

In [18]:

```
df=pd.read_csv(r"C:\Users\chila\Downloads\fiat500_VehicleSelection_Dataset.csv")
df
```

Out[18]:

	ID	model	engine_power	age_in_days	km	previous_owners	lat	lon	price
0	1	lounge	51	882	25000	1	44.907242	8.611560	8900
1	2	pop	51	1186	32500	1	45.666359	12.241890	8800
2	3	sport	74	4658	142228	1	45.503300	11.417840	4200
3	4	lounge	51	2739	160000	1	40.633171	17.634609	6000
4	5	pop	73	3074	106880	1	41.903221	12.495650	5700
...	...	...	...	...	...	...	...	...	...
1533	1534	sport	51	3712	115280	1	45.069679	7.704920	5200
1534	1535	lounge	74	3835	112000	1	45.845692	8.666870	4600
1535	1536	pop	51	2223	60457	1	45.481541	9.413480	7500
1536	1537	lounge	51	2557	80750	1	45.000702	7.682270	5990
1537	1538	pop	51	1766	54276	1	40.323410	17.568270	7900

1538 rows × 9 columns

In [19]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1538 entries, 0 to 1537
Data columns (total 9 columns):
#   Column          Non-Null Count  Dtype  
---  -
0   ID               1538 non-null   int64  
1   model            1538 non-null   object  
2   engine_power     1538 non-null   int64  
3   age_in_days      1538 non-null   int64  
4   km               1538 non-null   int64  
5   previous_owners  1538 non-null   int64  
6   lat              1538 non-null   float64 
7   lon              1538 non-null   float64 
8   price            1538 non-null   int64  
dtypes: float64(2), int64(6), object(1)
memory usage: 108.3+ KB
```

In [20]:

```
df.describe
```

Out[20]:

```
<bound method NDFrame.describe of
0      1 lounge      51      882  25000      1 \
1      2   pop      51     1186  32500      1
2      3 sport      74     4658 142228      1
3      4 lounge     51     2739 160000      1
4      5   pop      73     3074 106880      1
...    ...    ...    ...    ...    ...
1533 1534 sport     51     3712 115280      1
1534 1535 lounge    74     3835 112000      1
1535 1536   pop     51     2223  60457      1
1536 1537 lounge    51     2557  80750      1
1537 1538   pop     51     1766  54276      1

      lat      lon  price
0  44.907242  8.611560  8900
1  45.666359 12.241890  8800
2  45.503300 11.417840  4200
3  40.633171 17.634609  6000
4  41.903221 12.495650  5700
...    ...    ...    ...
1533 45.069679  7.704920  5200
1534 45.845692  8.666870  4600
1535 45.481541  9.413480  7500
1536 45.000702  7.682270  5990
1537 40.323410 17.568270  7900
```

[1538 rows x 9 columns]>

In [21]:

```
df.isnull().any()
```

Out[21]:

```
ID      False
model    False
engine_power  False
age_in_days  False
km        False
previous_owners  False
lat       False
lon       False
price     False
dtype: bool
```

In [23]:

```
df.isnull().sum()
```

Out[23]:

```
ID      0
model    0
engine_power  0
age_in_days  0
km        0
previous_owners  0
lat       0
lon       0
price     0
dtype: int64
```



In [24]:

```
df.loc[:10,["ID", "price"]]
```

Out[24]:

	ID	price
0	1	8900
1	2	8800
2	3	4200
3	4	6000
4	5	5700
5	6	7900
6	7	10750
7	8	9190
8	9	5600
9	10	6000
10	11	8950

In [26]:

```
df=df[["engine_power", "price"]]  
df.columns=["engine", "price"]
```

In [27]:

```
df.head()
```

Out[27]:


	engine	price
0	51	8900
1	51	8800
2	74	4200
3	51	6000
4	73	5700

In [30]:

```
sns.lmplot(x='engine', y='price', data=df, order=2, ci=None)
```

Out[30]:

<seaborn.axisgrid.FacetGrid at 0x246deb25480>



In [31]:

```
df.fillna(method='ffill')
```

Out[31]:

	engine	price
0	51	8900
1	51	8800
2	74	4200
3	51	6000
4	73	5700
...	...	...
1533	51	5200
1534	74	4600
1535	51	7500
1536	51	5990
1537	51	7900

1538 rows × 2 columns

In [32]:

```
x=np.array(df['engine']).reshape(-1,1)  
y=np.array(df['price']).reshape(-1,1)
```

In [34]:

```
df.dropna(inplace=True)
```

C:\Users\chila\AppData\Local\Temp\ipykernel\_10832\1552371080.py:1: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy) ([https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy))

```
df.dropna(inplace=True)
```

In [35]:

```
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.25)
```


In [37]:

```
regr=LinearRegression()  
regr.fit(x_train,y_train)  
print(regr.score(x_test,y_test))
```

0.02991002926555708

In [38]:

```
y_pred=regr.predict(x_test)  
plt.scatter(x_test,y_test,color='g')  
plt.plot(x_test, y_pred, color='k')  
plt.show()
```




In [41]:

```
df500=df[:][:500]  
sns.lmplot(x='engine',y='price', data=df500, order=1, ci=None)
```

Out[41]:


<seaborn.axisgrid.FacetGrid at 0x246df90a2f0>



In [43]:

```
df500.fillna(method='ffill',inplace=True)
x=np.array(df500['engine']).reshape(-1,1)
y=np.array(df500['price']).reshape(-1,1)
df500.dropna(inplace=True)
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.25)
regr=LinearRegression()
regr.fit(x_train,y_train)
print("Regression:",regr.score(x_test,y_test))
y_pred=regr.predict(x_test)
plt.scatter(x_test,y_test, color='y')
plt.plot(x_test, y_pred, color='k')
plt.show()
```

Regression: 0.02676853707836946



In [47]:

```
from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score
model=LinearRegression()
model.fit(x_train,y_train)
y_pred=model.predict(x_test)
r2=r2_score(y_test,y_pred)
print("R2 score:", r2)
```

R2 score: 0.02676853707836946


## Ridge regresssion and lasso

In [49]:

```
plt.figure(figsize=(10,10))
sns.heatmap(df.corr(),annot = True)
```

Out[49]:

<Axes: >



In [66]:

```
features= df.columns [0:2]
target = df.columns[-1]
#X and y values
x = df[features].values
y= df[target].values
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.3, random_state=7)
print("The dimension of X_train is {}".format(X_train.shape))
print("The dimension of X_test is {}".format(X_test.shape)) #Scale features
scaler = StandardScaler()
x_train = scaler.fit_transform(x_train)
x_test = scaler.transform(x_test)
```

The dimension of X\_train is (140, 2)

The dimension of X\_test is (60, 2)

In [65]:

```
lr = LinearRegression()
#Fit model
lr.fit(x_train, y_train)
#predict
#prediction = Lr.predict(X_test)
#actual
actual = y_test
train_score_lr = lr.score(x_train, y_train)
test_score_lr = lr.score (x_test, y_test)
print("\nLinear Regression Model:\n")
print("The train score for lr model is {}".format(train_score_lr))
print("The test score for lr model is {}".format(test_score_lr))
```

Linear Regression Model:

The train score for lr model is 0.08588880857157655  
The test score for lr model is 0.02676853707836946

In [70]:


```
ridgeReg = Ridge(alpha=10)
ridgeReg.fit(x_train,y_train)
#train and test scorefor ridge regression
train_score_ridge = ridgeReg.score(x_train, y_train)
test_score_ridge = ridgeReg.score (x_test, y_test)
print("\nRidge Model: \n")
print("The train score for ridge model is {}".format(train_score_ridge))
print("The test score for ridge model is {}".format(test_score_ridge))
```

Ridge Model:

The train score for ridge model is 0.9999080289793792  
The test score for ridge model is 0.9999063567964503

In [86]:

```
#plot size
plt.figure(figsize = (10, 10))
#add plot for ridge regression
plt.plot(features,ridgeReg.coef_,alpha=0.7,linestyle='none',marker='*',markersize=5,color='red',label=r'Ridge;  $\alpha=10$ ')
#add plot for lasso regression
plt.plot(lasso_cv.coef_,alpha=0.5,linestyle='none',marker='d',markersize=6,color='blue',label=r'lasso;  $\alpha = \text{grid}$ ')
#rotate axis
plt.xticks(rotation = 90)
plt.legend()
plt.title("Comparison plot of Ridge, Lasso and Linear regression model")
plt.show()
```





In [76]:

```
#Lasso regression model
print("\nLasso Model: \n")
lasso = Lasso (alpha = 10)
lasso.fit(x_train,y_train)
train_score_ls =lasso.score(x_train,y_train)
test_score_ls =lasso.score (x_test,y_test)
print("The train score for ls model is {}".format(train_score_ls))
print("The test score for ls model is {}".format(test_score_ls))
```

Lasso Model:

The train score for ls model is 0.9999717830403652

The test score for ls model is 0.9999716845769463

**conclusion: The above model is fit for linear regression.**